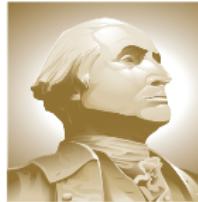


Google RECSIMs

Evaluating Recommender Systems

Hema Puchakayala (Group 12)



THE GEORGE
WASHINGTON
UNIVERSITY
WASHINGTON, DC

- Project Description
- Outline
- Markov Decision Process (MDP) Formulation
- Code
- Results
- Outlook

Project Description

- Evaluate how effectively recommender systems respond to the evolving interests of users over time
- Recommend over 1100+ of news categories to users
- Synthetic users and documents are generated by sampling features

Outline

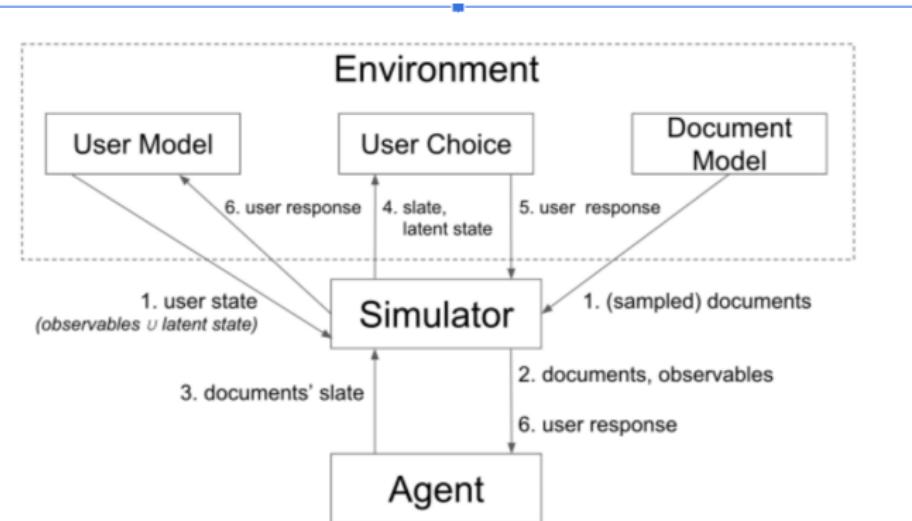


Figure 2: Control flow (single user) in the RecSim architecture.

- **DriftingDocument**

Holds documents d_i with features: x topic, p popularity score, and q quality score:

$$d_i = (x_i, p_i, q_i)$$

- **DriftingUserState**

Stores user's state $u_{j,t}$ at time t : latent preferences $\theta_{j,t}$, satisfaction $\sigma_{j,t}$, fatigue $\phi_{j,t}$, timestep $\tau_{j,t}$:

$$w_{j,t} = (\theta_{j,t}, \sigma_{j,t}, \phi_{j,t}, \tau_{j,t})$$

- **States (S):**

User state at time t : $s_t = (\theta_{j,t}, \sigma_{j,t}, \phi_{j,t})$

- **Actions (A):**

Choice of document(s) to recommend at each step

- **Transition Model (P):**

Defined by the user state update equations:

$$\theta_{j,t+1} = (1 - \alpha)\theta_{j,t} + \alpha x_{i,t} + \epsilon_{j,t}$$

$$\sigma_{j,t+1} = (\alpha^2)\sigma_{j,t} + (1 - \alpha)^2 p_i + 2\alpha(1 - \alpha)q_i + \epsilon_\sigma$$

$$\phi_{j,t+1} = (\alpha^2)\phi_{j,t} + 2\alpha(1 - \alpha)p_i + (1 - \alpha)^2 q_i + \epsilon_\phi$$

- **Reward Function (R):**

User click or engagement, as modeled by the response equation:

$$\Pr(\text{click}_{j,i,t} = 1) = \sigma(\theta_{j,t}^T x_i + \beta_1 q_i + \beta_2 p_i + \gamma_1 \sigma_{j,t} + \gamma_2 \phi_{j,t} + \xi_{j,i,t})$$

- **DriftingResponseModel**

Maps user/document features into click probabilities:

$$\Pr(\text{click}_{j,i,t} = 1) = \sigma(\theta_{j,t}^T x_i + \beta_1 q_i + \beta_2 p_i + \gamma_1 \sigma_{j,t} + \gamma_2 \phi_{j,t} + \xi_{j,i,t})$$

- **DriftingUserModel**

Advances user state by updating preferences with drift:

$$\theta_{j,t+1} = (1 - \alpha)\theta_{j,t} + \alpha x_{i,t} + \epsilon_{j,t}$$

$$\sigma_{j,t+1} = (\alpha^2)\sigma_{j,t} + (1 - \alpha)^2 p_i + 2\alpha(1 - \alpha) q_i + \epsilon_\sigma$$

$$\phi_{j,t+1} = (\alpha^2)\phi_{j,t} + 2\alpha(1 - \alpha) p_i + (1 - \alpha)^2 q_i + \epsilon_\phi$$

Environment Setup

```
class DriftingEnvironment:  
    """Combines all components into a full environment."""  
    def __init__(self, num_users, num_documents, alpha, a, b, beta_1, beta_2, gammal, gamma2, categories)  
        self.num_users = num_users  
        self.num_documents = num_documents  
        self.categories = categories  
        self.alpha = alpha  
        self.a = a  
        self.b = b  
        self.beta_1 = beta_1  
        self.beta_2 = beta_2  
        self.rng = np.random.default_rng(seed)  
  
        self.doc_sampler = DriftingDocumentSampler(categories, alpha=alpha, a=a, b=b, seed=seed)  
        self.user_sampler = DriftingUserSampler(categories, alpha=alpha, a=a, b=b, seed=seed)  
        self.response_model = DriftingResponseModel(beta1=beta_1, beta2=beta_2, gammal = gammal,gamma2=gamma2, seed=seed)  
        self.user_model = DriftingUserModel(alpha=0.01, seed=seed)  
  
    def step(self, user, doc):  
        response = self.response_model.simulate_response(user, doc)  
        user = self.user_model.update_state(user, doc)  
        return int(response), user  
  
    def reset(self):  
        users = self.user_sampler.sample_users(self.num_users)  
        documents = self.doc_sampler.sample_documents(self.num_documents)  
        return users, documents
```

```
for step in range(NUM_ROUNDS):
    step_reward = 0.0
    for user_index, user in enumerate(users):

        click_probs = np.array([env.response_model.score(user, doc) for doc in documents])
        state_key = tuple(np.round(click_probs, 4))

        action = model.esoft(state_key, len(documents))
        selected_document = documents[action]

        reward, updated_user = env.step(user, selected_document)
        tdq_cum_rew += reward
        users[user_index] = updated_user

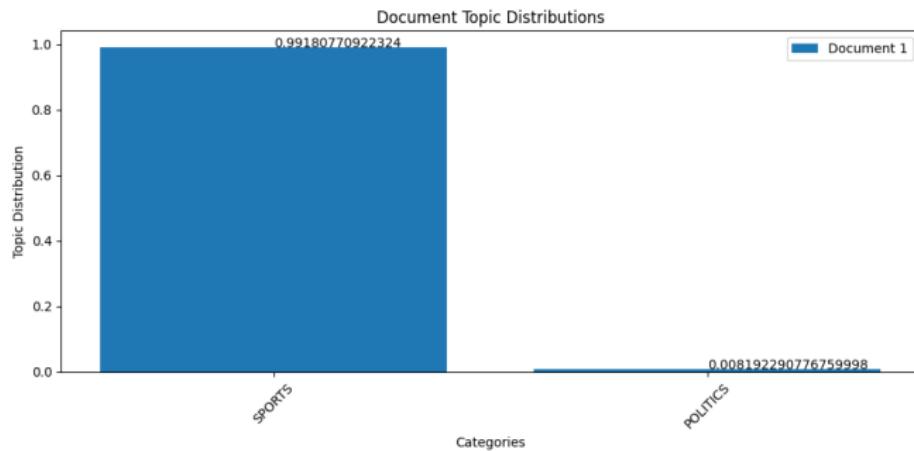
        scalar_reward = reward if np.isscalar(reward) else np.sum(reward)

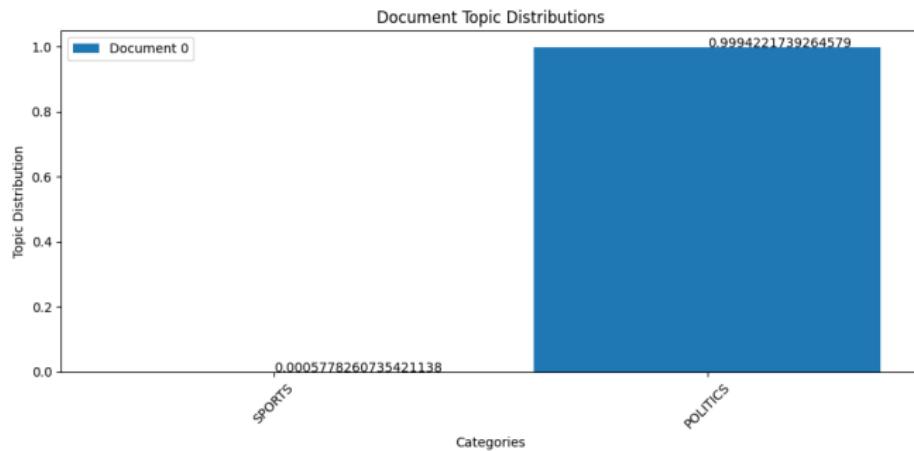
        latent_preference_list_tdq.append(updated_user.theta.copy())
        sigma_list_tdq.append(updated_user.sigma)
        phi_list_tdq.append(updated_user.phi)

        click_probs_next = np.array([env.response_model.score(updated_user, doc) for doc in documents])
        next_state_key = tuple(np.round(click_probs_next, 4))
        model.update(action, state_key, next_state_key, scalar_reward)
```

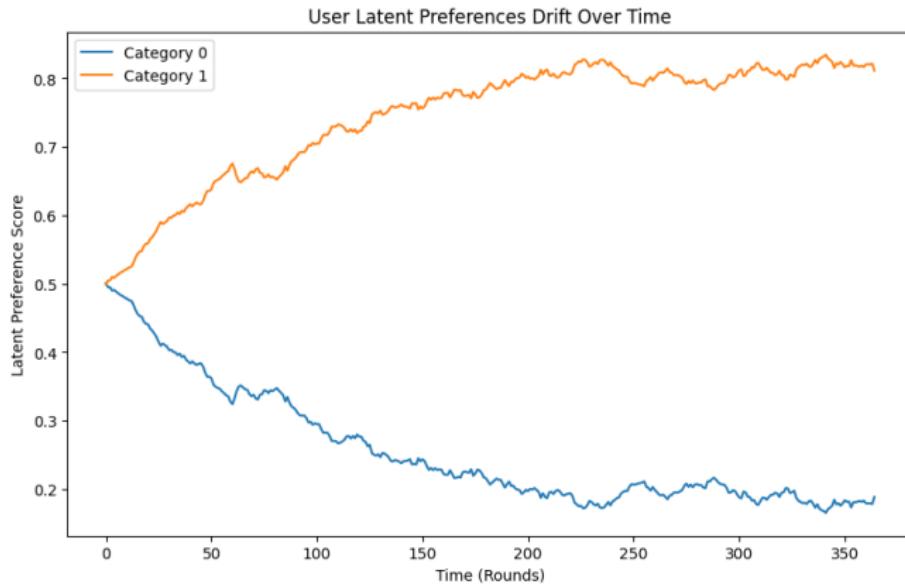
- Compared the performance of both a random model and a Q-learning model for recommending documents to users

Sports



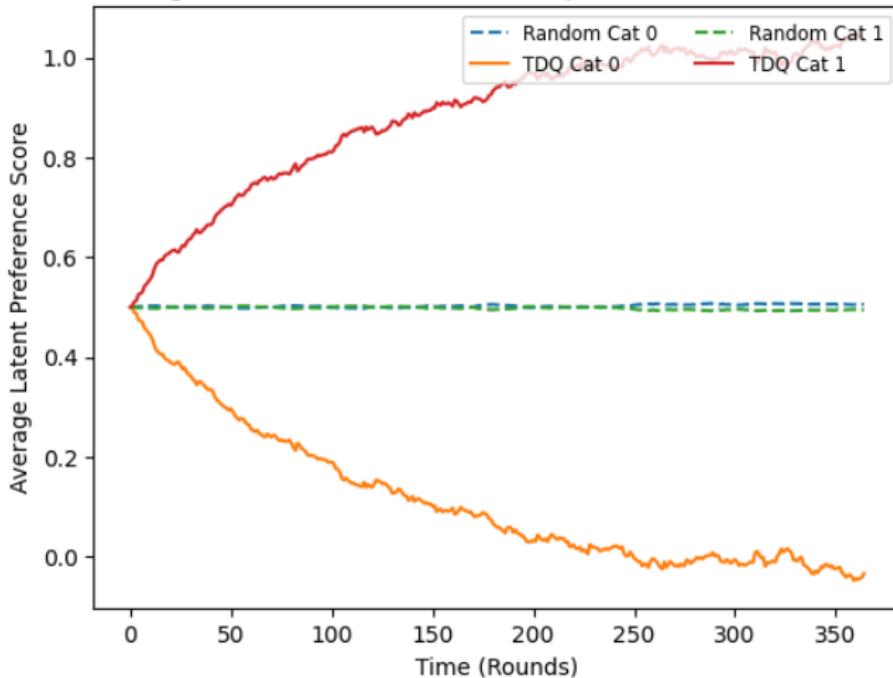


- Actual Latent Preference Drift



- Latent Preference Drift by the model

Average Latent Preferences Over Episodes: Random vs TDQ



- The current model is a simulation for studying adaptive recommender systems
- Compare the Random Model and Q-Learning Model
- Scaling up for larger user base and document categories